

"WHAT I WON'T BUILD" IN NLP

BY RACHEL TATMAN @ ALP 2020 CONF.

WHAT'S MOST URGENT "DOGMA" IN NLP?

- ① Used to be: "Language technologies must work equally well for everyone."
- ② Now is: "Language technologists [NLP researchers & developers] must proactively mitigate harm from our work."

↳ eg minimizing differences in accuracy between subgroups

↳ eg "should this be built at all?" - we have limited resources & need to use them

wisely. We don't need to ^{immediately} research everything just because it's there.

Just like particle physics

KEY QUESTION

What do we mean by "responsible"?

"Do you consider yourself responsible for the usages imagined from the applications / algorithms you create?" (Ethical grounds)

STANDARD ETHICS THAT WE NLP RESEARCHERS SHOULD ABIDE BY (ACL ACM CODE OF ETHICS)

- ① contribute to society & to human well-being, acknowledging that all people are stakeholders in computing
- ② Avoid harm
- ③ Be honest & trustworthy
- ④ Be fair & take action not to discriminate
- ⑤ Respect the work req'd to produce new ideas, inventions, creative works, & computing artifacts
- ⑥ Respect privacy

politics screws all of this up, and it should, sometimes.

whose?

at what cost? ←

⑦ Honor confidentiality.

WHAT SHE WON'T BUILD

→ deceptive tech that tries to pass itself as a human (without telling the human it's interacting with)

ies in nicely
to linguistics
of textok. ch.

→ social category detectors → using sociolinguistic identity to predict your race, gender...

→ surveillance technology

QUESTIONS TO ASK

① Who benefits from this system existing?

— harmed?

② Can users choose NOT to interact with this system?

③ Does this system re-inforce systemic inequalities?

④ Does this genuinely better the world?